

Toward Impactful SE Research by Taking Roads Less Traveled

Zhendong Su

University of California, Davis



What is the **key mission** of
Computer Science?

To help people **turn creative**
ideas into working systems

Software research, especially **SE**, is
at the center of this mission





Two instances


❑ Validate production compilers

- ◆ Black-box analysis



❑ Analyze floating-point software

- ◆ Dynamic analysis




Validate Production Compilers



black box

Compiler complexity



LLVM bug 14972

```
struct tiny { char c; char d; char e; };

void foo(struct tiny x) {
    if (x.c != 1) abort();
    if (x.e != 1) abort();
}

int main() {
    struct tiny s;
    s.c = 1; s.d = 1; s.e = 1;
    foo(s);
    return 0;
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```


Developer comment

“... very, very concerning when I got to the root cause, and very annoying to fix ...”

http://llvm.org/bugs/show_bug.cgi?id=14972

Vision

P \equiv



Key challenges

□ Generation

- ◆ How to generate **different**, yet **equivalent** tests?

□ Validation


- ◆ How to check that tests are **indeed equivalent**?

□ Both are long-standing hard issues



Equiv. modulo inputs

- Relax equiv. wrt a **given input i**
 - ◆ Must: $P(i) = P_k(i)$ on input i
 - ◆ Okay: $P(j) \neq P_k(j)$ on all input $j \neq i$
- Exploit close **interplay** between
 - ◆ **Dynamic** program execution on **some input**
 - ◆ **Static** compilation for **all input**


Profile




Mutate




Mutate





....



equivalent wrt 1



Find bugs



$O' \neq O$



Revisit challenges

□ Generation (**easy**)

- ◆ How to generate **different**, yet **equivalent** tests?

□ Validation (**easy**)

- ◆ How to check that tests are **indeed equivalent**?

□ ~~Both are long-standing hard issues~~

LLVM bug 14972

```
struct tiny { char c; char d; char e; };

void foo(struct tiny x) {
    if (x.c != 1) abort();
    if (x.e != 1) abort();
}

int main() {
    struct tiny s;
    s.c = 1; s.d = 1; s.e = 1;
    foo(s);
    return 0;
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```

Seed file

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) abort();
    if (x.d != 20) abort();
    if (x.e != 30) abort();
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) abort();
    if (z.c != 12) abort();
    if (z.d != 22) abort();
    if (z.e != 32) abort();
    if (l != 123) abort();
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
```

Seed file

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) abort();
    if (x.d != 20) abort();
    if (x.e != 30) abort();
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) abort();
    if (z.c != 12) abort();
    if (z.d != 22) abort();
    if (z.e != 32) abort();
    if (l != 123) abort();
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

← unexecuted

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
```

Transformed file

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) /* deleted */;
    if (x.d != 20) abort();
    if (x.e != 30) /* deleted */;
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) /* deleted */;
    if (z.c != 12) abort();
    if (z.d != 22) /* deleted */;
    if (z.e != 32) abort();
    if (l != 123) /* deleted */;
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```

Reduced file

```
struct tiny { char c; char d; char e; };

void foo(struct tiny x) {
    if (x.c != 1) abort();
    if (x.e != 1) abort();
}

int main() {
    struct tiny s;
    s.c = 1; s.d = 1; s.e = 1;
    foo(s);
    return 0;
}
```


```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```


LLVM bug autopsy

```
struct tiny { char c; char d; char e; };
```

```
void foo(struct tiny x) {  
    if (x.c != 1) abort();  
    if (x.e != 1) abort();  
}
```

GVN: load struct
using 32-bit load



```
int main() {  
    struct tiny s;  
    s.c = 1; s.d = 1; s.e = 1;  
    foo(s);  
    return 0;  
}
```

```
$ clang -m32 -O0 test.c ; ./a.out  
$ clang -m32 -O1 test.c ; ./a.out  
Aborted (core dumped)
```


LLVM bug autopsy

```
struct tiny { char c; char d; char e; };
```

```
void foo(struct tiny x) {  
    if (x.c != 1) abort();  
    if (x.e != 1) abort();  
}
```


```
int main() {  
    struct tiny s;  
    s.c = 1; s.d = 1; s.e = 1;  
    foo(s);  
    return 0;  
}
```

GVN: load struct
using 32-bit load



SRoA: read past
the struct's end

→
undefined
behavior




```
$ clang -m32 -O0 test.c ; ./a.out  
$ clang -m32 -O1 test.c ; ./a.out  
Aborted (core dumped)
```

LLVM bug autopsy

```
struct tiny { char c; char d; char e; };
```


```
void foo(struct tiny x) {  
    if (x.c != 1) abort();  
    if (x.e != 1) abort();  
}
```

GVN: load struct
using 32-bit load




```
int main() {  
    struct tiny s;  
    s.c = 1; s.d = 1; s.e = 1;  
    foo(s);  
    return 0;  
}
```

SRoA: read past
the struct's end



→
undefined
behavior



remove



```
$ clang -m32 -O0 test.c ; ./a.out  
$ clang -m32 -O1 test.c ; ./a.out  
Aborted (core dumped)
```

Seed file

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) abort();
    if (x.d != 20) abort();
    if (x.e != 30) abort();
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) abort();
    if (z.c != 12) abort();
    if (z.d != 22) abort();
    if (z.e != 32) abort();
    if (l != 123) abort();
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
```

Transformed file

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) /* deleted */;
    if (x.d != 20) abort();
    if (x.e != 30) /* deleted */;
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) /* deleted */;
    if (z.c != 12) abort();
    if (z.d != 22) /* deleted */;
    if (z.e != 32) abort();
    if (l != 123) /* deleted */;
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```

GCC bug 58731

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        for (; c;)
            for (;;) {
                e = a > 2147483647 - b;
                if (d) break;
            }
    return 0;
}
```

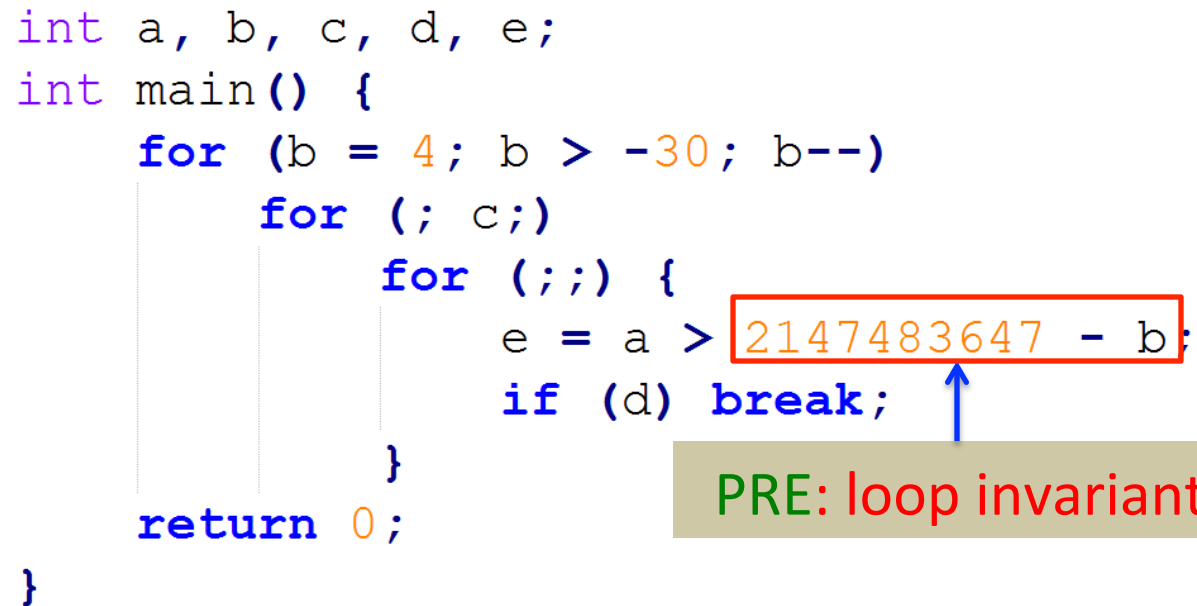
```
$ gcc -O0 test.c ; ./a.out
```

```
$ gcc -O3 test.c ; ./a.out
```

```
^C
```

GCC bug autopsy

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        for (; c;)
            for (;;) {
                e = a > 2147483647 - b;
                if (d) break;
            }
    return 0;
}
```




The diagram illustrates the nested loop structure of the code. Vertical lines connect the opening and closing braces of the three nested loops: the outermost loop (for b), the middle loop (for c), and the innermost loop (for ;). A red box highlights the constant value 2147483647 in the expression `e = a > 2147483647 - b;`. A blue arrow points from a text box below to this constant.

PRE: loop invariant

```
$ gcc -O0 test.c ; ./a.out
$ gcc -O3 test.c ; ./a.out
^C
```

GCC bug autopsy

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        int f = 2147483647 - b;
        for (; c;)
            for (;;) {
                e = a > f;
                if (d) break;
            }
    return 0;
}
```



A diagram illustrating a GCC bug. A green box labeled "LIM" has a blue arrow pointing to the constant value "2147483647" in the code snippet. The constant is highlighted with a red box. The variable "f" is also highlighted with a red box in the line "e = a > f;".

```
$ gcc -O0 test.c ; ./a.out
```

```
$ gcc -O3 test.c ; ./a.out
```

```
^C
```


GCC bug autopsy

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        int f = 2147483647 - b;
        for (; c;)
            for (;;) {
                e = a > f;
                if (d) break;
            }
    return 0;
}
```

integer overflow

```
$ gcc -O0 test.c ; ./a.out
$ gcc -O3 test.c ; ./a.out
^C
```

Seed program



```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        for (; c;)
            for (;;) {
                b++;
                e = a > 2147483647 - b;
                if (d) break;
            }
    return 0;
}
```

no longer a
loop invariant

```
$ gcc -O0 test.c ; ./a.out
$ gcc -O3 test.c ; ./a.out
```


Orion (PLDI'14)

Prune dead code




Athena (OOPSLA'15)

Prune & inject dead code



Hermes (OOPSLA'16)

Mutate live code



Bug counts

	GCC	LLVM	TOTAL
Reported	710	594	1304
Fixed	467	309	776

- **ISSTA'15**: Stress-testing link-time optimization
- **ICSE'16**: Analyzing compilers' diagnostic support
- **PLDI'17**: Skeletal program enumeration

LLVM 3.9 & 4.0 Release Notes

“... thanks to **Zhendong Su and his team**
whose fuzz testing **prevented many bugs**
going into the release ...”

EMI

❑ General, widely applicable

- ◆ Can test compilers, analysis, transformation tools
- ◆ Finds real-world tests and requires no reference compilers


❑ Very effective realizations

- ◆ Uncovered **1300+ bugs** in GCC and LLVM in 3 years
- ◆ Many **miscompilations** and **long latent** bugs

	GCC	LLVM	TOTAL
Reported	710	594	1304
Fixed	467	309	776

***Take-away:** Not only do good
research, but be its **loyal**,
continuous user!*

Analyze Floating-Point Software



Floating-point code



- **Important:** bugs can lead to disasters
- **Challenging:** hard to get right

Why difficult?

- **FP Math** \neq **Real Math**
- **Non-linear** relations
- **Transcendental** functions
 \sin , \log , \exp , ...

```
1  double foo(double x){  
2      if (x<=1.0)  
3          x++;  
4  
5      y = x*x;  
6      if (y<=4.0)  
7          x--;  
8      return x;  
9  }
```

Challenging for all known approaches

New perspective: ME

(p, ϕ)

Analyzing numerical programs

- Coverage-based testing
- Boundary value analysis
- Numerical exception detection

Floating-point constraint solving

↓ Mathematical
Execution (ME)

r Mathematical optimization (MO)

input x drives p to satisfy $\phi \iff x$ minimizes r

FP constraints

Solving the floating-point constraint π

$$(SIN(x) == x) \wedge (x \geq 10^{-10})$$


- ▶ Satisfiable if x is floating-point

$$\text{For } x \in \mathbb{F}, SIN(x) = x \Leftrightarrow x \simeq 0$$

- ▶ Unsatisfiable if x is real

$$\text{For } x \in \mathbb{R}, SIN(x) = x \Leftrightarrow x = 0$$


Step 1



Simulate π with a floating-point program R

- ▶ $R(x) \geq 0$ for all x
- ▶ $R(x) = 0 \Leftrightarrow x \models \pi$

Step 2



Minimize R as if it is a mathematical function

- Let x^* be the minimum point

$$\pi \text{ satisfiable} \Leftrightarrow R(x^*) = 0$$

Construct R

Necessary Conditions to meet :

1. $R(x) \geq 0$ for all x
2. $R(x) = 0 \Leftrightarrow x \models \pi$

How?


Constraint π	Program R
$x == y$	$(x - y)^2$
$x \leq y$	$x \leq y ? 0 : (x - y)^2$
$\pi_1 \wedge \pi_2$	$R_1 + R_2$
$\pi_1 \vee \pi_2$	$R_1 * R_2$

R can be constructed from a CNF form

Minimize R

Unconstrained programming techniques:

- ▶ Local optimization
- ▶ Monte Carlo Markov Chain (MCMC)
- ▶ We use them as black-box
- ▶ Do not analyze π ; execute R



Theoretical guarantees

Let R satisfy (1) $R(x) \geq 0$, and (2) $R(x) = 0 \Leftrightarrow x \models \pi$, and x^* be a minimum point of R . Then

$$\pi \text{ satisfiable} \Leftrightarrow R(x^*) = 0.$$

Threats

- ▶ Floating-point inaccuracy when calculating with R
- ▶ Sub-optimal x^*

Example

$$(SIN(x) == x) \wedge (x \geq 10^{-10})$$



$$(SIN(x) - x)^2 + \begin{cases} 0 & \text{if } x \geq 10^{-10} \\ (x - 10^{-10})^2 & \text{otherwise} \end{cases}$$



$$x^* = 9.0 * 10^{-9} \text{ (can be others)}$$

XSat & results

- Developed the ME-based XSat tool
- Evaluated against **MathSat** and **Z3**
- Used **SMT-Comp 2015 FP** benchmarks
- Result summary
 - **100%** consistent results
 - **700+X** faster than MathSat
 - **800+X** faster than Z3

***Take-away:** Don't be afraid of
difficult problems, look at them from
new perspectives,
even for “damned” problems!*

Generalizations

- ❑ Coverage-based testing of FP code
- ❑ Boundary value analysis
- ❑ FP exception detection
- ❑ Path divergence detection

ME in the long run

- ❑ Offers a new general analysis paradigm
- ❑ Complements existing approaches
 - ◆ Random concrete execution (CE)
 - ◆ Symbolic execution (SE)
 - ◆ Abstract execution (AE)

Additional Reflections

Pursue “**less-crowded**” directions

Much work on **applying** ML/DL/AI
but much less on

- ❑ **Formalizing** ML/DL/AI systems
- ❑ **Testing/verifying/explaining** them
- ❑ **Making** them more **usable** & **effective**

Pursue “**less-crowded**” directions

- ❑ **Intelligent programming assistants**
- ❑ **Compilers for developers** (not machines)
- ❑ **Knowledge capture & reuse**
- ❑ Anything **leading closer to key mission**

Many exciting
opportunities,
challenges,
uncertainties